

MEASUREMENT // RESTRAINT // ADVERSARIAL REVIEW

Measuring exposed AI without becoming the threat

How LeakyCompute separates evidence, limits its claims, and used Daybreak Blue to red team its own safety boundary.

**ACTIVE
PROBING
PAUSED**

Mahdi Hedhli // 26 August 2026

Three measurements, never summed

Each number answers a different question and carries a different confidence level.

19,348

archive listings

Captured 27 July 2026.
Historical and not reverified.

19,376

indexed records

Public records as of 25 August.
Counted, not probed.

1,360


reverified hosts

Retained read-only
observations within 180 days.

Current index coverage is entirely Shodan-derived. Source bias and blind spots are part of the result.

Rules before features

A host must pass every gate before a historical verification request could leave the runner.



GET only • no redirects • 32 KiB body cap • reachability, never demonstrated impact

Three surfaces, three trust boundaries

The smallest useful representation crosses each boundary.

PUBLIC

Aggregate counts

No raw addresses

RESEARCH LAB

Minimised corpus

Verified Access identity +
allowlist

ADMIN

Ingest ✦ authorization

Separate token and exact
routes

A signed identity is not a display name. A researcher route is not an admin route. An aggregate is not a raw record.

What open endpoints make reachable

Host and stack pairs in the retained historical snapshot. A host can appear in more than one class.

969 models and inference open

Ollama, Open WebUI, LocalAI, compatible APIs




188 workflow execution open

ComfyUI workflow endpoints



129 training metadata open

MLflow and TensorBoard



78 code execution reachable

Jupyter and Ray prerequisites



Reachability observed • exploitability not tested

Daybreak Blue attacked the safety claims

The review covered code, identity, storage, workflows, Cloudflare configuration, runners, and tests. No third-party target was probed.

IDENTITY

- researcher impersonation
- alias lifecycle
- denial-log privacy

PROBES

- address bypasses
- stale provenance
- ledger crash window

STATE

- partial recounts
- fail-open exclusions
- route confusion

AUTOMATION

- public raw artifacts
- unsafe event parsing
- supply-chain drift

Every promise became a question: how could this fail under hostile input, stale state, or a partial crash?

Fix the bug, contain the capability

Concrete defects were remediated. Risky surfaces stayed off when the architectural control was not ready.

FIXED OR CONTAINED

- signed identity verification
- canonical public-IP parsing
- exact bridge routes
- redacted workflows
- bounded reads + redirects blocked
- pinned supply chain

STILL REQUIRED

- durable pre-probe lease
- pageable corpus storage
- verifiable opt-out deletion
- retention at corpus scale
- complete reconciliation
- durable security state

13 discovery artifact and workflow log sets were deleted. No credential leak was confirmed.

A narrower system, and an honest one

The safety boundary includes identity, logs, retries, storage, redirects, workflows, and dependency pins. The probe is only one piece.

AVAILABLE

Passive index aggregates
Approved researcher lab
Local owner-operated checks

SUSPENDED

Scheduled active discovery
Hosted target checks
Public override endpoints

leakycompute.mahdihedhli.com